

A Diagnostic No-Selection Report for Train-Free LLM Anchoring Mitigation

Anonymous ACL submission

Abstract

LLM anchoring-bias work often reports prompt-level mitigation gains without carrying the full integrity burden of pre-registered controls, parser failures, cost, and held-out confirmation. We evaluate a train-free anchoring-mitigation screen on source-backed numerical questions about publication year and author count, using a hosted gpt-5-mini backend with temperature 0, a 64-token response cap, and a strict single-number parser. The frozen dev/validation screen contains 893 terminal model-call records across 51 minimal-pair groups, six method conditions, two target quantities, and predeclared metrics for anchor pull, mitigation lift, query cost, parser failures, and no-anchor template stability. The main validation diagnostics show measurable anchor displacement and several positive lift values, including 0.55353484 for ignore-anchor instructions and 0.63428962 for reference-validity triage on author-count targets, but all candidates fail the selection contract because no-anchor stability is unavailable, preservation and corrected held-out confirmation are absent, and parser/API failures remain material. The supported conclusion is therefore negative: no mitigation method is selected, and the paper contributes an auditable no-selection decision rather than a positive debiasing claim. We therefore scope the manuscript as a diagnostic report, not as a final cross-benchmark or held-out mitigation paper.

1 Introduction

Anchoring is a classic judgment bias in which an initial number can shift later estimates (Tversky and Kahneman, 1974; Epley and Gilovich, 2006; Strack and Mussweiler, 1997). The same behavioral risk matters for LLMs because numerical prompts and weakly relevant references can shift answers even when the target quantity should be recoverable from context (Lou and Sun, 2024; Huang et al., 2025). Prompt-sensitivity work makes this

harder to evaluate: a prompt intervention can look useful on one slice while merely changing parser success, verbosity, or sensitivity to prompt wording (Razavi et al., 2025; Taubenfeld et al., 2025). Mechanistic localization work gives a separate motivation for future white-box tests, but behavior-only screens cannot identify pathways or circuits by themselves (Owusu et al., 2026).

This creates an evidence gap for train-free mitigation claims. A reviewer should not accept a method because it has a high validation lift if the same evaluation lacks no-anchor stability controls, preservation checks for informative references, parser/API failure accounting, cost accounting, multiple-testing discipline, or final held-out confirmation. The practical question is therefore not merely whether a mitigation produces a better-looking validation number, but whether it survives the full evidence contract: pre-test mechanism, fixed protocol, no look-ahead, complete trial logging, parser accountability, cost accounting, preservation of informative references, multiple-testing discipline, and final held-out confirmation.

This draft reports a negative result under that contract. The reported evidence consists only of dev/validation diagnostics over source-backed publication-year and author-count questions. The screen used a hosted gpt-5-mini backend at temperature 0 with a 64-token response cap and a strict parser that accepted exactly one finite number. It produced 893 terminal model-call records: 876 successes, 12 parser failures, and 5 API errors. The measured validation diagnostics include positive author-count lift for ignore-anchor instructions (0.55353484) and reference-validity triage (0.63428962), but these values do not select a mitigation because no-anchor template stability is unavailable, preservation checks and corrected held-out confirmation are missing, and failures are part of the evidence. Our contribution is a compact integrity case for no selection: the screen states the

pre-test mechanism, fixes the protocol and costs before testing, reports the dev/validation evidence with its limits, and refuses to promote a high diagnostic metric into a method claim.

The negative result is useful because it separates three questions that are often conflated. First, did the model move numerically when an anchor was present? The validation rows say yes for the author-count target: the no-mitigation comparator has absolute anchor displacement 0.32275132. Second, did a train-free prompt intervention reduce that displacement on the same diagnostic slice? Some interventions did: reference-validity triage has validation author-count lift 0.63428962, ignore-anchor instructions have 0.55353484, and salient-number quarantine has 0.55109290. Third, is any intervention therefore selected as a reliable mitigation? No. The third question requires controls that the present screen does not have. This paper is organized around that distinction so that the positive-looking diagnostic values do not become an unsupported method claim.

The review-stage target is therefore not to maximize a backtest score or to choose the visually best row. It is to make the no-selection decision auditable. Every reported number is either a row-level diagnostic value, a run-level call/failure count, or an availability count from the frozen analysis tables. The paper avoids a test/OOS headline, does not report no-anchor stability, and does not infer internal mechanisms. Those absences are not cosmetic caveats; they are the reason the selection decision is negative.

2 Related Work

Classic anchoring accounts emphasize insufficient adjustment and selective accessibility (Tversky and Kahneman, 1974; Epley and Gilovich, 2006; Strack and Mussweiler, 1997). In those accounts, an anchor can shape the range of candidate answers before the subject has fully evaluated whether the anchor is relevant. That mechanism translates naturally to LLM prompting: a number in the prompt can become a salient token sequence, a plausible reference point, or a spurious constraint on the answer distribution. A train-free mitigation should therefore do more than suppress any number; it should reduce irrelevant-anchor pull while preserving genuinely informative numerical context.

LLM-specific anchoring studies show that irrelevant or weakly relevant numerical cues can shift

model answers, while mitigation results remain sensitive to prompt design and evaluation protocol (Lou and Sun, 2024; Huang et al., 2025; Valencia-Clavijo, 2025). This sensitivity is the central reason that the present screen treats parser failures, cost, and control availability as part of the result. If a prompt intervention changes answer format or parser success, a nominal reduction in anchor pull may be a formatting artifact. If an intervention asks for multiple independent estimates, a better diagnostic value may partly reflect higher query budget rather than a cleaner debiasing mechanism. If the no-anchor control has only one template, template sensitivity cannot be separated from anchor sensitivity.

Prompt-sensitivity and uncertainty work motivates treating apparent mitigation gains cautiously when controls are incomplete (Razavi et al., 2025; Taubenfeld et al., 2025; Lyu et al., 2025). The present paper follows that caution by requiring a predeclared gate rather than selecting the row with the largest validation lift. Mechanistic localization work motivates future white-box analysis, but behavior-only screens cannot identify pathways or circuits by themselves (Owusu et al., 2026). For that reason, the mechanism in this paper is behavioral and economic: the screen evaluates whether a low-cost, train-free intervention can reduce measured anchor pull under fixed cost and parser constraints. It does not claim to locate the internal computation that produced the response.

3 Benchmark and Protocol

The evaluation protocol and metric definitions were fixed before testing. Dev rows were not used for headline claims. Validation rows are used only for diagnostic selection replay, and no test rows are used in this draft. The metric map includes signed anchor pull, absolute anchor displacement, mitigation lift, query-cost overhead, and no-anchor template stability. The current evidence makes no-anchor stability unavailable because each group has only one no-anchor template.

The benchmark is intentionally narrow and inspectable. Each item is a source-backed paper-metadata question whose target is either publication year or author count. These targets were chosen because the correct answer is numeric, the prompt can contain a potentially distracting numerical anchor, and a strict parser can determine whether the model returned exactly one finite num-




Figure 1: Conceptual overview of the diagnostic anchoring screen. Source-backed numeric prompts are evaluated for anchor pull through train-free procedures, but selection is blocked by the integrity gate when controls are missing. The figure is a structural schematic; the numerical evidence remains in the ledger-grounded tables and diagnostic plots.

ber. The screen is not a broad benchmark suite. It is a diagnostic source family designed to expose whether candidate prompt procedures survive the integrity floor before any held-out or cross-source claim is attempted.

The frozen run contains 51 valid primary minimal-pair groups and 893 terminal model-call records. The run ledger records 876 successful calls, 12 parser failures, and 5 API errors. Those failures remain in the analysis surface. They are not dropped as nuisance rows because a train-free mitigation that improves a numeric metric while making answer extraction brittle would not be a reliable mitigation. The parser contract is therefore part of the evaluation: the model had to emit one finite number under a 64-token cap, and failures were counted against the evidence package.

Metric availability is also part of the protocol. The analysis table has 140 method/target/split metric rows, split evenly across dev and validation. Of these rows, 92 are available, 24 are call-status summaries, and 24 are unavailable by protocol limit. Every unavailable row is the no-anchor stability metric under the one-template no-anchor design. This matters because no-anchor stability is the con-

trol that would distinguish anchor-specific improvement from generic prompt instability. Without that control, a validation lift can be reported as a diagnostic value but cannot certify a selected mitigation.

Selection was frozen as a gate rather than a ranking exercise. A candidate needed validation-only evidence that survived parser quality, cost, preservation, corrected selection, and held-out confirmation requirements. The combination replay records no selected method. This replay is not a post-hoc attempt to rationalize a weak result; it is the implementation of the pre-test GO/NO-GO rule. The paper therefore reports validation diagnostics, failure accounting, and no-selection logic, but withholds any test/OOS headline.

Appendix Figure 6 summarizes the evidence-handling flow from fixed protocol to screened calls, metric audit, selection replay, and no-go decision.

4 Train-Free Methods

The screen compares a no-mitigation comparator, two baselines, and three candidate train-free interventions. The no-mitigation condition is retained

Setup field	Paper-facing value
Benchmark/source	Source-backed paper-metadata questions; one diagnostic source family
Task count/split	51 minimal-pair groups; 893 terminal dev/validation calls; held-out test split quarantined
Model/backend	hosted gpt-5-mini
Budget/decoding	temperature 0, top-p 1, one response, max 64 output tokens
Parser	strict single finite-number extraction with failures counted
Methods/baselines	no mitigation; ignore-anchor instruction; reflection/chain-of-thought; reference-validity triage; blind independent re-estimation; salient-number quarantine
Metrics	anchor pull, mitigation lift, query-cost overhead, parser/API failures, no-anchor template stability
Selection rule	validation-only selection replay with parser, cost, preservation, corrected-selection, and held-out confirmation gates

Table 1: Paper-facing setup facts for the train-free anchoring screen.

as the primary comparator because it estimates how strongly the answer moves under the anchored prompt without an explicit mitigation. It is not a mitigation and is not eligible for positive selection.

The blanket ignore-anchor instruction is a baseline rather than a proposed final method. Its behavioral mechanism is simple: tell the model to disregard unrelated anchors and answer directly. This can reduce pull when the anchor is plainly irrelevant, but it can also encourage the model to discard useful numerical context if relevance is ambiguous. The method uses one call per item, so its query-cost overhead is zero in the current accounting.

Reflection/checking is included as a diagnostic reasoning baseline. It asks the model to check its reasoning before returning a final numeric answer. The plausible mechanism is that deliberation may weaken the first salient number’s influence by forcing a consistency check against the question. The corresponding risk is that additional reasoning text can change output format or introduce a new salient number; the strict parser therefore sees only the final numeric field.

Reference-validity triage is the clearest candidate mechanism in the screen. It asks the model to classify prompt numbers as irrelevant, weakly informative, informative, or diagnostic before estimating the answer. Its intended effect is selective:




Figure 2: Metric availability blocks positive selection: no-anchor stability is unavailable under the current one-template control design.

downweight irrelevant anchors while preserving references that genuinely describe the target. This is the kind of mechanism a reviewer could accept in principle, but the present evidence still does not select it because the required stability and held-out controls are absent.

Blind independent re-estimation attacks anchoring by repetition and aggregation. It uses three anchor-quarantined calls and combines independent estimates under frozen median or confidence rules. The behavioral mechanism is reduced dependence on one anchored context, but the economic cost is explicit: validation query-cost overhead is 2.00000000 for both author-count and publication-year targets. A higher-cost method cannot be selected merely because it has access to more calls.

Salient-number quarantine masks or moves irrelevant numerical anchors while preserving task semantics and diagnostic references. Its mechanism is to reduce visual and lexical salience of irrelevant numbers before estimation. This is attractive when the anchor is clearly unrelated, but it depends on the preprocessing step not damaging the target information. In the current screen it remains a diagnostic candidate, not a selected mitigation.

All candidates are rejected as main mitigations in the current draft because the selection gates are not satisfied. This rejection is not based on low lift alone. It is based on the combination of unavailable no-anchor stability, absent preservation checks, material parser/API failures, query-cost accounting, missing corrected held-out confirmation, and the explicit no-selection replay.

Condition	Calls/item	Operational role
No mitigation	1	comparator using the frozen condition prompt
Ignore-anchor instruction	1	baseline directive to ignore unrelated anchors
Checking	1	diagnostic reasoning check before final numeric answer
Reference-validity triage	1	classifies numbers by relevance before estimation
Blind independent re-estimation	3	combines anchor-quarantined independent estimates
Salient-number quarantine	1	masks or moves irrelevant anchors while preserving task semantics

Table 2: Train-free method procedures and query budgets used in the diagnostic screen.




Figure 3: Validation absolute anchor displacement is diagnostic only and is not test/OOS evidence.

5 Results

5.1 Anchor Displacement

The validation anchor-pull diagnostics show that the screen produced measurable behavior under the frozen metric definitions, but these values do not select a method. The author-count target is the clearest diagnostic slice. The no-mitigation comparator has validation absolute anchor displacement 0.32275132. Reference-validity triage has 0.11803351, salient-number quarantine has 0.14488536, ignore-anchor instructions have 0.14409722, reflection/checking has 0.28333333, and blind independent re-estimation has 0.29867725. These values are directionally informative: several interventions reduce author-count displacement relative to no mitigation. They are not sufficient for a main mitigation claim because the selection rule was not “choose the small-




Figure 4: Validation mitigation lift includes positive and negative diagnostic values, but no candidate is selected because integrity gates fail.

est validation displacement.”

Publication-year rows behave differently. The corresponding validation absolute displacements are near zero for the comparator and several interventions. That pattern is itself a warning against overgeneralizing from a single target quantity. A method that appears useful on author count may have little meaningful room to improve on publication year because the baseline displacement is already extremely small. The diagnostic matrix therefore reports source, target, backend, metric, and scale together rather than collapsing the screen into one aggregate score.

5.2 Lift Is Not Selection

Validation lift makes the temptation to overclaim visible. On author-count rows, reference-validity triage has lift 0.63428962, ignore-anchor instructions have 0.55353484, salient-number quarantine has 0.55109290, reflection/checking has 0.12213115, and blind independent re-estimation has 0.07459016. A metric-only selector would likely prefer triage or ignore-anchor instructions on this slice. The predeclared selector does not. It asks whether a method has enough evidence to be accepted as a mitigation after parser, stability, cost, preservation, corrected validation, and held-out confirmation gates.

The publication-year lift rows reinforce the same point from the other direction. Several publication-year validation lift values are negative, including -1.00000000 for salient-number quarantine and -2.00049407 for both blind re-estimation and reference-validity triage. Because publication-year displacement is already close to zero in the comparator, small absolute changes can produce unstable relative lift. The paper therefore treats lift as a diagnostic lens rather than a headline outcome.

Benchmark/source	Task count/split	Model/backend	Metric	Key takeaway
Paper metadata screen	893 dev/validation calls	gpt-5-mini, temp 0	call status	876 successes; 12 parser and 5 API failures
Paper metadata screen	validation author count	gpt-5-mini, temp 0	anchor pull	no-mitigation absolute pull is 0.32275132
Paper metadata screen	validation author count	gpt-5-mini, temp 0	lift	ignore-anchor lift is 0.55353484 but insufficient alone
Paper metadata screen	validation author count	gpt-5-mini, temp 0	lift	triage lift is 0.63428962 but insufficient alone
Paper metadata screen	validation author count/year	gpt-5-mini, temp 0	cost/stability	blind re-estimation costs 2.00000000 extra queries; no-anchor stability unavailable

Table 3: The consolidated diagnostic matrix exposes source family, scale, backend, metrics, and numerical takeaways; because it is a single-source validation screen, it supports no selected mitigation.




Figure 5: Parser/API failures are material diagnostic evidence against unqualified selection.

This is why the abstract reports positive author-count lift but immediately states that all candidates fail the selection contract.

5.3 Failure Accounting

Parser and API failures are part of the evaluation record rather than cleanup noise. Across the frozen ledgered screen, the failure stratification records 12 parser-failure rows and 5 API-error rows. This matters because the evaluated interventions are prompt procedures. A prompt procedure can change not only the estimate, but also whether the answer is extractable under the strict single-number parser. If the method’s apparent improvement depends on rows where failed parses disappear, the result is not an anchoring mitigation result. The strict parser and failure accounting prevent that failure mode.

The call-status evidence also limits how much the reader should infer from positive rows. The 893 terminal model calls include 876 successes, so the screen is not dominated by failures. At the same time, the 17 non-success terminal records are material because the selection decision is binary. A selected mitigation would need to be reliable enough for the evaluator to recommend it under

fixed protocol and fixed cost. The present failure surface is compatible with diagnostic analysis, but not with promotion to a main method.

5.4 Cost and Search Breadth

Cost accounting gives a second reason not to rank by one metric. Most methods use one call per item, so their validation query-cost overhead is 0.00000000. Blind independent re-estimation uses three calls and has validation cost overhead 2.00000000 for both author-count and publication-year targets. Its mechanism is plausible, but it buys a different amount of computation. Because the protocol treats query budget as part of the intervention, the extra calls cannot be ignored when deciding whether a method is selected.

The search breadth is modest but fully disclosed. The screen covers six method conditions: no mitigation, ignore-anchor instructions, reflection/checking, reference-validity triage, blind independent re-estimation, and salient-number quarantine. It covers two target quantities, dev and validation splits, and six metric families when call status is counted. The metric-availability audit records 24 call-status summary rows, 24 signed-anchor-pull rows, 24 absolute-anchor-pull rows, 24 no-anchor-stability rows, 24 query-cost rows, and 20 mitigation-lift rows. This breadth is enough to audit the diagnostic decision; it is not enough to claim a cross-benchmark mitigation.

5.5 Threshold Replay

The threshold replay ties these pieces together; Appendix Figure 7 gives the compact visual version. The combination replay records ‘no method selected’ after applying the validation-only selection logic. The no-selection result is not a claim

that every method is useless. It is a claim that the current evidence package is not strong enough to certify any method as a reliable mitigation. The strongest diagnostic values remain useful for designing the next benchmark round: triage and quarantine reduce author-count displacement with one call, while blind re-estimation illustrates the cost of repeated independent estimates. None of those observations changes the current decision.

The most important missing control is no-anchor stability. The metric-availability audit has 24 rows unavailable by protocol limit, all from no-anchor stability under the one-template control design. Without multiple no-anchor templates, the evaluator cannot tell whether an intervention reduces anchor pull specifically or merely changes sensitivity to the unanchored prompt wording. This is why no-anchor stability is treated as a gate rather than an optional robustness statistic.

The final missing piece is held-out confirmation. The draft uses dev and validation diagnostics only. It deliberately has no test/OOS headline. That makes the report conservative, but it also means the current manuscript cannot honestly satisfy a long-paper empirical standard that requires cross-source and held-out evidence. The review-stage conclusion is therefore narrow: the screen supports an auditable no-selection decision, not a selected mitigation and not a general benchmark claim.

6 Mechanism Discussion

The behavior observed in validation diagnostics is consistent with anchoring or prompt sensitivity under the frozen protocol. It is not evidence for an attention pathway, activation circuit, or localization result. The candidate methods make different behavioral bets: ignore-anchor instructions test whether a direct directive suppresses irrelevant anchors; reference-validity triage asks the model to classify numerical relevance before estimation; blind independent re-estimation reduces dependence on one anchored prompt at a measured two-query validation overhead; and salient-number quarantine reduces surface salience while trying to preserve task semantics.

These mechanisms are interpretable, but none is certified. A high lift value without no-anchor stability could be prompt-template sensitivity; without preservation evidence it could be over-suppression of useful references; with parser failures it could be an extraction artifact; with extra calls it could be a

budget artifact; and without held-out confirmation it could be validation selection. The current evidence helps prioritize future evaluation, especially triage and quarantine, but it does not authorize deployment of any method as a mitigation.

7 Implications for the Next Evaluation

The next evaluation should be built around the observed blockers rather than the largest validation lift. It should add multiple no-anchor templates per group, preservation tests for informative and diagnostic numerical references, independent source families beyond paper metadata, and one held-out confirmation run after freezing the candidate and thresholds. A future positive claim should report the full searched method set, failed candidates, parser/API failures, query costs, and corrected selection logic alongside any headline.

8 Traceability and Selection Audit

The no-selection decision is only useful if a reviewer can reconstruct it from quantitative evidence rather than trusting prose alone. The supplementary evidence table therefore ties each reported decision statement to the corresponding analysis record, split scope, and support status. Anchor displacement and mitigation lift are reported only where the fixed metric definition is available. Parser/API failures are reported as failure counts, and no-anchor stability is reported as unavailable rather than imputed. This is why the limitation is treated as an observed protocol fact rather than as missing exposition.

This traceability also keeps the validation-only boundary visible. The paper-facing evidence for anchor displacement, mitigation lift, query-cost overhead, and selection replay is scoped to dev/validation diagnostics, and the manuscript does not use the quarantined held-out split for the reported no-selection conclusion. That distinction matters because the strongest positive-looking values are validation author-count values. They can motivate another benchmark round, but they cannot become an out-of-sample headline after the fact.

The selection replay is deliberately stricter than a metric leaderboard. It combines the observed lift values with query cost and unavailable no-anchor stability before producing the no-method-selected decision. This prevents three common failure modes: choosing the largest validation lift after searching multiple methods, ignoring extra

calls for multi-query procedures, or treating an unavailable control metric as if it had passed. Under that replay, reference-validity triage and salient-number quarantine remain plausible candidates for future study, but neither is promoted to a selected mitigation.

Failure accounting is handled the same way. The paper reports the 12 parser failures and 5 API errors because the evaluated objects are prompt procedures, not offline formulas. A method that changes output format can change the strict parser’s success rate; a method that relies on extra attempts can change the cost surface; and an API failure can remove evidence from a cell. The current screen is successful enough to be analyzed, with 876 successful terminal calls, but the non-success records remain part of the integrity argument.

The conceptual figures are treated as schematics rather than evidence. Their role is to orient the reader to the diagnostic screen and evidence-handling flow, while the numerical claims remain attached to tables, plots, and supplementary analysis records. The scientific conclusion therefore does not depend on a visual impression: the no-selection decision comes from the screened calls, metric availability audit, failure accounting, cost comparison, and selection replay.

9 Conclusion

The evidence-grounded conclusion is no method selected. The six evaluated conditions are useful diagnostics, but none is accepted as a main mitigation under the current evidence package. Validation lift identifies promising directions, especially triage and quarantine, but selection requires multiple no-anchor templates, preservation metrics, independent source families, corrected validation selection, and a final held-out run under the same ledgered protocol.

10 Limitations

The draft has four hard limitations. First, it has no positive mitigation selection. Second, it has no test/OOS headline evidence. Third, no-anchor template stability is unavailable for all current method/target/split cells. Fourth, behavior-only data cannot support attention/pathway claims. These limitations determine what the evidence can support.

The most important limitation is evidence scope. The current benchmark is one diagnostic source

family: source-backed paper-metadata questions. It is useful because the target answers are numeric and auditable, but it is not a cross-benchmark suite. A long-paper empirical claim would need independent source families and a readable cross-source matrix. It would also need held-out/OOS headline evidence after the selection rule is frozen. The present draft does not have that evidence and therefore does not claim it.

The second limitation is the one-template no-anchor control. The no-anchor stability metric is unavailable in 24 method/target/split rows because the current manifest has only one no-anchor template per group. This is a design limitation, not a missing calculation. Without additional no-anchor templates, selection remains blocked even when validation lift is positive.

The third limitation is preservation. Some interventions reduce irrelevant-anchor pull by telling the model to ignore numbers or by masking salient numbers, but the current screen does not contain a reportable preservation metric strong enough to certify that useful numerical references are retained.

The fourth limitation is failure accounting and evidence traceability. The screen has 12 parser-failure rows and 5 API-error rows, and the present figures are diagnostic summaries rather than standalone proof. The claims table and source figures remain necessary because they bind each numerical statement to analysis rows and ledgered records.

11 Ethical Considerations

The main ethical risk in this study is overstatement. A prompt procedure that appears to reduce anchoring on one validation slice could be adopted as a general debiasing recipe even when stability, preservation, parser reliability, and held-out confirmation are missing. The manuscript therefore reports positive diagnostic values only inside the no-selection frame and avoids deployment language. This is especially important for numerical decision support, where an apparently harmless prompt instruction could change estimates while hiding sensitivity to irrelevant numbers.

The evaluation also treats cost and failure transparency as responsible-use requirements. Multicall procedures consume more inference budget, and strict parsing failures can remove difficult cases from a naive metric. Reporting query overhead, parser failures, and API errors reduces the chance that a method is selected because it silently changes

extraction behavior or spends a different amount of computation. The same principle motivates the held-out split quarantine: a final test/OOS claim should be made only after the candidate and thresholds are frozen.

The benchmark uses source-backed paper-metadata questions rather than private user data or personal attributes. Even so, the results should not be generalized to safety-critical numerical advice. The present evidence supports an audit of a diagnostic screen, not a guarantee that train-free prompting can prevent anchoring in deployed systems.

References

- Nicholas Epley and Thomas Gilovich. 2006. [The anchoring-and-adjustment heuristic](#).
- Yiming Huang, Biquan Bie, Zuqiu Na, Weilin Ruan, Songxin Lei, Yutao Yue, and Xinlei He. 2025. [Understanding the anchoring effect of llm with synthetic data: Existence, mechanism, and potential mitigations](#). arXiv preprint.
- Jiaxu Lou and Yifan Sun. 2024. [Anchoring bias in large language models: An experimental study](#). arXiv preprint.
- Yougang Lyu, Shijie Ren, Yue Feng, Zihan Wang, Zhumin Chen, Zhaochun Ren, and Maarten de Rijke. 2025. [Self-adaptive cognitive debiasing for large language models in decision-making](#). arXiv preprint.
- Hillary N. Owusu, Sarah Wiegrefe, and Naomi H. Feldman. 2026. [Localizing anchoring pathways in language models](#). arXiv preprint.
- Amirhossein Razavi, Mina Soltangheis, Negar Arabzadeh, Sara Salamat, Morteza Zihayat, and Ebrahim Bagheri. 2025. [Benchmarking prompt sensitivity in large language models](#). arXiv preprint.
- Fritz Strack and Thomas Mussweiler. 1997. [Explaining the enigmatic anchoring effect: Mechanisms of selective accessibility](#).
- Amir Taubenfeld, Tom Sheffer, Eran Ofek, Amir Feder, Ariel Goldstein, Zorik Gekhman, and Gal Yona. 2025. [Confidence improves self-consistency in llms](#). arXiv preprint.
- Amos Tversky and Daniel Kahneman. 1974. [Judgment under uncertainty: Heuristics and biases](#).
- Felipe Valencia-Clavijo. 2025. [Anchors in the machine: Behavioral and attributional evidence of anchoring bias in llms](#). arXiv preprint.

Anchoring Evidence Pipeline




Figure 6: Conceptual evidence-handling pipeline for the train-free anchoring diagnostics. The pre-test rule fixes the mechanism and GO/NO-GO contract; minimal pairs are screened under the fixed protocol; metric audit and selection replay precede the no-go decision. The pipeline summarizes evidence handling and does not assert a selected mitigation.




Figure 7: Threshold replay rejects every baseline/candidate as a main mitigation under current evidence.

A Evidence Trace

The supplementary evidence table links each quantitative decision statement to exact diagnostic records. These records support auditability of the frozen dev/validation evidence and are not test/OOS evidence.

B Regeneration

A deterministic replay rebuilds the report tables, figures, and evidence table from the frozen protocol records. The replay interface is limited to the evidence bundle and does not govern stage decisions.